

Mean Inference

Estimating a population average

Download the section 13 .Rmd handout to
STAT240/lecture/sect13-single-mean.

Material in this section is covered by Chapter 13 on
the notes website.

Download the file TIM.txt to STAT240/data.

The Boston Marathon is a prestigious annual 26.2 mile race.

It is held in April, but was held in October in 2011.

TIM.txt contains times of Boston Marathon runners from 2010 to 2011.

Possible questions of interest:

- What is the average running time μ ?
- Is there a difference between μ_{2010} and μ_{2011} ?

Analysis is based on the CLT.

We assume the observations, X_i are i.i.d. with mean μ and sd σ .

$$X_i \sim D(\mu, \sigma)$$

Let's focus on the 18-34 female finishers from 2010.

The point estimate for μ is the sample mean

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$$

and we observed $\bar{x} = 235.5$.

Let's look at a simulation for \bar{X} .

The CLT gives

$$\bar{X} \sim N\left(\mu, \frac{\sigma}{\sqrt{n}}\right)$$

for large enough n , which will be the basis of our CI and H-test.

A CI, in general, is

point estimate \pm critical value \times standard error

- For μ , the point estimate is \bar{x}
- The standard error is $\frac{\sigma}{\sqrt{n}}$

We should be able to use a Z critical value:

$$\frac{\bar{X} - \mu}{\sigma/\sqrt{n}} \sim N(0, 1)$$

Problem: we don't have σ . The quantity

$$\frac{\bar{X} - \mu}{S/\sqrt{n}}$$

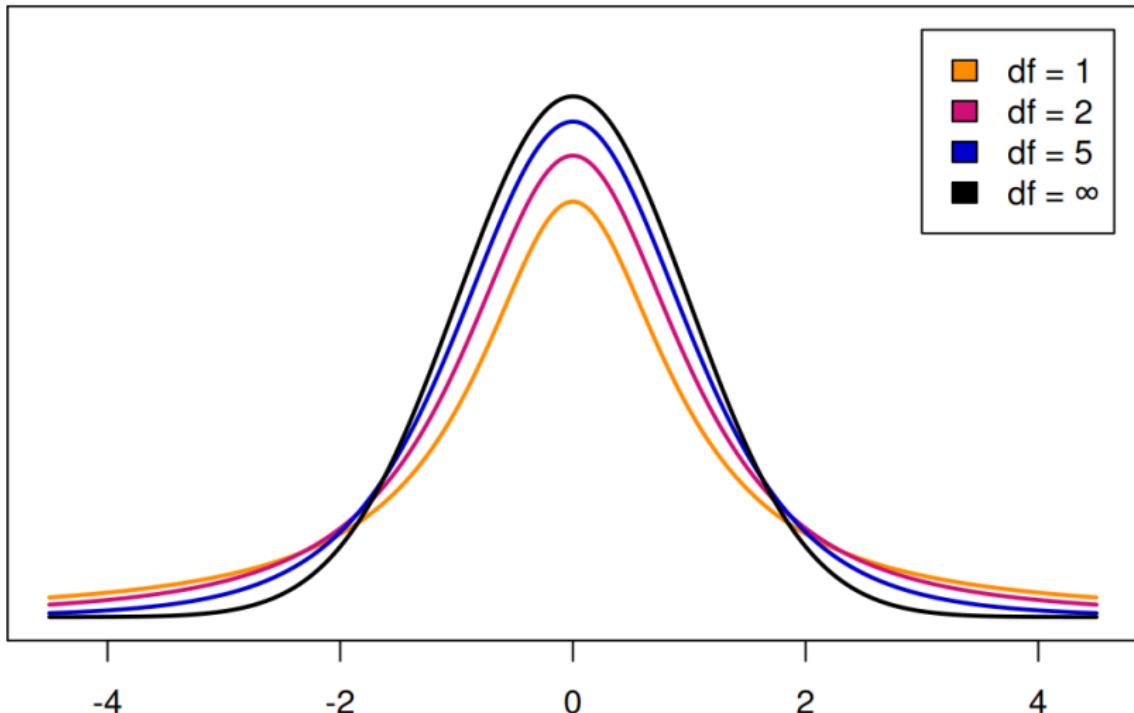
is not *quite* normal.

The sampling distribution for \bar{X} when σ is unknown uses the **Student's T distribution**.

$$\frac{\bar{X} - \mu}{S/\sqrt{n}} \sim T$$

This is similar to $N(0, 1)$, but wider.

T distributions

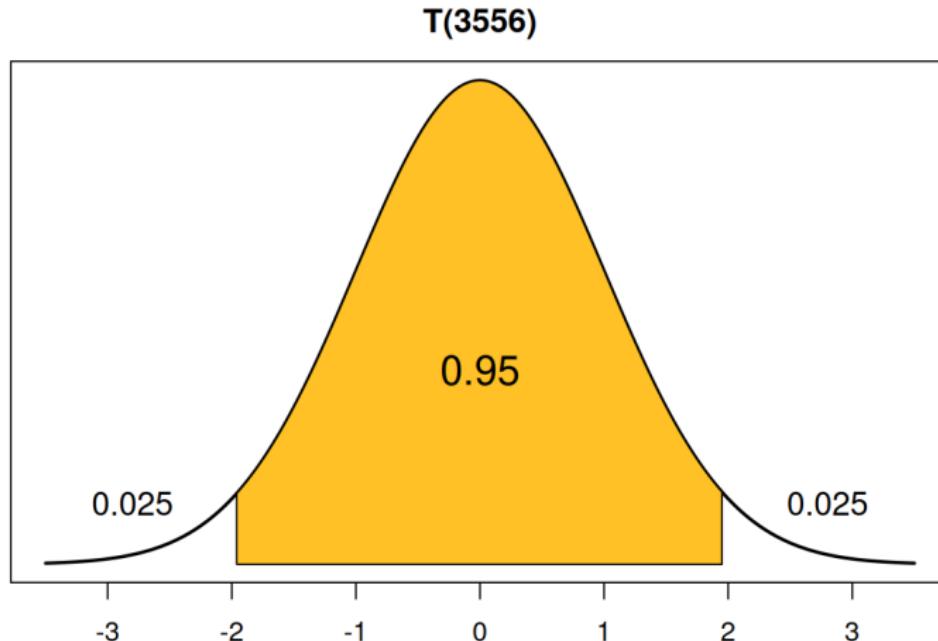


The T has heavier tails than $N(0, 1)$, controlled by degrees of freedom.

In single mean inference, $df = n - 1$.

Find critical values (quantiles) with `qt()`.

For 95% confidence:



T CI for μ :

$$\bar{X} \pm t_{\alpha/2, n-1} \cdot \frac{s}{\sqrt{n}}$$

We are 95% confident that the true average time for an 18-34 year old female runner in 2010 is within (234.3, 236.7).

Check with `t.test()`.

Is the true average time for an 18-34 year old female runner in 2010 equal to 240 minutes?

$$H_0 : \mu = 240 \quad \text{versus} \quad H_A : \mu \neq 240$$

We gather evidence against H_0 with $\bar{x} = 235.5$.

Let's use $\alpha = 0.05$.

If $\mu = 240$, then the quantity

$$\frac{\bar{X} - 240}{\sigma / \sqrt{n}}$$

should be $N(0, 1)$, implying a Z test.

But, we don't know σ .

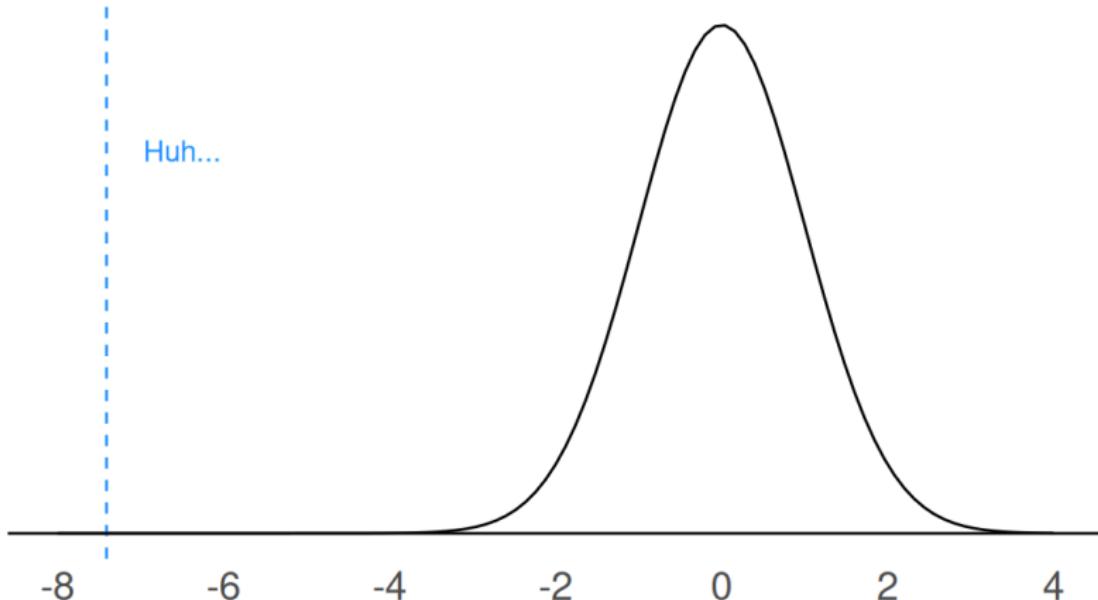
Instead, use

$$\frac{\bar{X} - 240}{S/\sqrt{n}} \sim T_{n-1}$$

Our observed test statistic is -7.41 . Is this consistent with our null T distribution?

$T(3556)$

Huh...



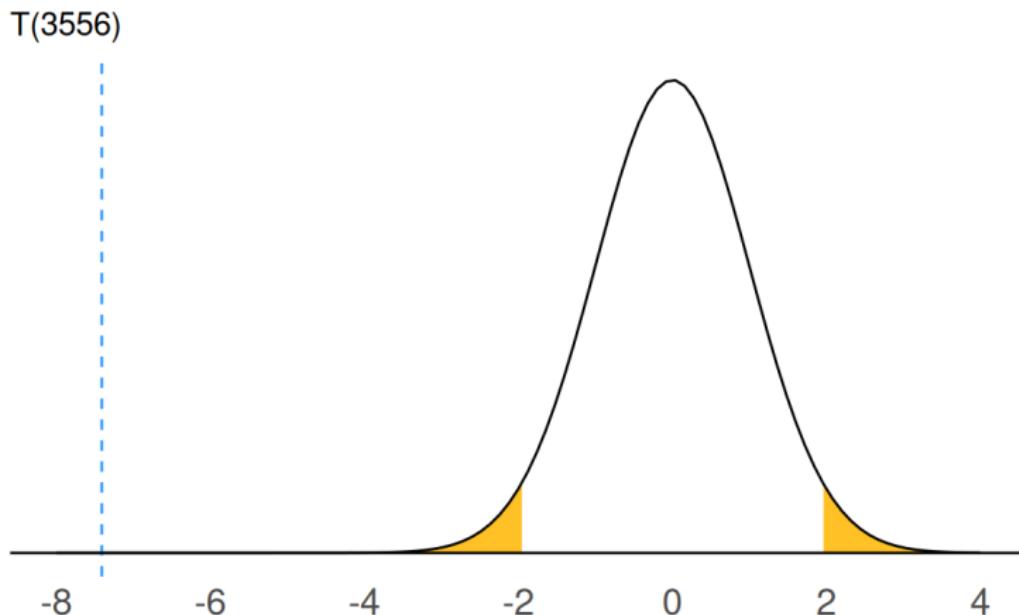
Our p-value

$$P(T_{n-1} \leq -7.41) + P(T_{n-1} \geq 7.41)$$

is very small, so we reject H_0 .

We can check this result in `t.test()`.

A rejection region also rejects H_0 .



Recall that our 95% CI for μ did not contain 240: (234.3, 236.7).

Similarly, we rejected the null value 240 at the 5% level. In fact, the methods are equivalent.

The following 2 statements are equivalent for T CIs and hypothesis tests for μ :

- A $100(1 - \alpha)\%$ CI contains μ_0 .
- A test for $H_A : \mu \neq \mu_0$ at level α rejects H_0 .

Conversely, if μ_0 is inside of the $100(1 - \alpha)\%$ CI, we would fail to reject H_0 in a two-sided test.

Consider female runners age 60-64 from 2010. Do we have evidence that their average time is *less* than 280 minutes?

- Write one-sided hypotheses about μ
- Re-calculate summary statistics
- Perform the T test with $\alpha = 0.02$.
- Check with `t.test()`

In general, perform a T test on independent draws from the same population. Use hypotheses

$$H_0 : \mu = \mu_0 \quad \text{versus} \quad H_A : \mu \neq \mu_0$$

and test statistic

$$T = \frac{\bar{X} - \mu_0}{S/\sqrt{n}}$$

and null distribution T_{n-1} .

P-value calculation:

- If the alternative is $H_A : \mu < \mu_0$, the p-value is the area *below* the test statistic.
- If the alternative is $H_A : \mu > \mu_0$, the p-value is the area *above* the test statistic.
- If the alternative is $H_A : \mu \neq \mu_0$, the p-value is $2 \times$ the area *outside* of the test statistic.